



Know what changed, whether it's true, and **why it matters to you.**

An autonomous, evidence-backed analyst that watches competitors, markets and regulators — for strategy, product-marketing and compliance teams.

Industry domain **Business**



Keeping up with competitors and regulators is still a manual, weekly grind



Checked by hand

Dozens of competitor, news and regulator pages, re-checked every week against a spreadsheet.



Alerts add noise

They report that *something* changed — often the same story many times — not whether it matters.



Nothing remembers

Evidence lives in screenshots, history in one person's head. When they are busy, nobody is watching.

THE REAL QUESTION

- ① What materially changed?
- ② Is it actually true?
- ③ Why does it matter to us?
- ④ Who needs to know?

Agents can now do an analyst’s legwork. Trust still needs evidence and memory.

WHAT’S NEW

AI agents can plan research, use tools – search, open pages, read documents – and check their own work.

Until recently, that legwork needed a person.


WHAT’S MISSING FOR BUSINESS USE

	A GENERAL AI ASSISTANT	WHAT TRUST REQUIRES
Works	– When you ask	✓ Continuously, on a schedule
Memory	– Starts from zero	✓ A dated record of every fact
Proof	– Links	✓ Quotes checked word for word
Context	– Generic	✓ Your company and your teams



And the economics work — if AI is used sparingly: plain code for most of the work, models only where judgment is needed.

One loop: set up once with you, then it runs on its own



 **World state** — a dated, versioned, evidence-backed memory of every entity and fact. Every step reads from it and writes to it.

 Agent loop  Single model call  Rules in code  You decide

Facts with proof. The agent’s assessment, labelled.

Razorpay / Pricing & fees / Pricing change

🕒 Historical · recorded, never alerted

Razorpay’s pricing page adds “0%* platform fees for first 90 days”

EVIDENCE STATUS

SEVERITY

TEAMS

✔ Confirmed

📊 High

Product · Sales & BD · Strategy

What changed

Facts

BEFORE · 9 MAR 2026 CAPTURE

No “0%*” platform-fee offer on the page

AFTER · 6 JUL 2026 CAPTURE · STILL LIVE

“0%* platform fees for first 90 days”

✔ Quote verified

on razorpay.com/pricing, archived and live · primary source · supports

Why it matters to Orbit Payments

Agent assessment

Illustrative

A 90-day fee holiday lowers the cost of switching gateways for small merchants — Orbit’s core segment. Expect price objections in new-merchant deals while the offer runs.

Assumes: the offer applies to new merchants only — its terms were not verified.

A real change, found on day one

Replaying archived captures of razorpay.com/pricing shows when the offer appeared:

- 9 Mar 2026 Offer absent
- × Apr 2026 Empty JavaScript shell, skipped
- × Jun 2026 Empty JavaScript shell, skipped
- ◆ 6 Jul 2026 Offer present

4 of 10 monthly captures since Sep 2025 were empty shells. The quality gate skipped them instead of reporting “everything removed”.

Real: page text, quotes and capture dates, checked 28 Sep 2026. Illustrative: severity, teams and “why it matters”, written for a fictional competitor, Orbit Payments. The card is a mock-up of the product’s layout.

How sure are we? Rules decide – not the model.

FIVE EVIDENCE STATUSES, COMPUTED BY CODE

 Confirmed	A primary source supports it: the company's own site, the issuing regulator or a statutory registry.
 Corroborated	No primary source, but two or more independent publishers agree and none contradict.
 Single source	Exactly one independent publisher. Severity is capped below Critical.
 Conflicting	Primary sources disagree or contradict the claim — or independent reports clash with no primary source to settle it.
 Unverified	No verified supporting evidence yet. Severity is capped at Medium.

Evaluated in a fixed order, conflict first; the first match wins. Syndicated copies count once; forums never move the status.



Quotes checked word for word

Evidence counts only if its quote is found in a page the agent actually opened. Changed numbers fail.



Facts kept apart from interpretation

The assessment is labelled as the agent's and lists its assumptions.



Web pages can't give orders


Fetches text is treated as data. Research tools are read-only.



Humans approve anything external


Sharing outside the company, external email and webhooks wait for a person.

The cheapest test goes first — most checks never reach a model




Filtered, not forgotten: every removed change is listed with its tier and the reason.

Two agent loops; everything else is deterministic



Custom agent runtime — runs, steps, typed tools, budgets, guardrails, traces

Each step is one validated decision. Budgets cap steps, tool calls, time and cost; web text is fenced as untrusted; every call is traced with tokens and cost.



Built end to end – and plain about what has been verified

385

automated backend tests pass

- ✓ Unit tests for the rules: evidence statuses, severity caps, clustering, scheduling, quote checks, SSRF, robots.txt
- ✓ Full-loop integration on PostgreSQL, with a scripted model and a simulated web: request > plan > approval > baseline > change > filter > verify > card > route > feedback

razorpay.com

real-world checks, 28 Sep 2026

- ✓ Live pricing page extracted with the 2% headline, the 0%* offer and the 2.15% + GST footnote intact
- ✓ robots.txt read and honoured
- ✓ SSRF guard refuses the cloud-metadata address
- ✓ Web-archive history: 10 captures, 4 empty shells rejected, the offer change located

Every screen

of the web app, built

- ✓ Onboarding, live planner and plan review
- ✓ Dashboard, intelligence feed and cards
- ✓ World state, monitoring, learned rules
- ✓ Run traces, approvals and demo lab
- ✓ Build, type-check and lint pass; clicked through against the real backend



Live runs need your own API keys — one model provider (Anthropic, OpenAI or Gemini) and one search provider (Tavily, Exa or Serper). So far, the model-driven steps have run against a scripted model in tests.

Beachhead: B2B teams without a CI team

IDEAL CUSTOMER PROFILE				
Hypothesis to validate with design partners				
COMPANIES	MARKETS	COMPETITORS	TEAM	FIRST MARKET
B2B, roughly 100–2,000 employees	Fast-moving and public: payments, fintech, B2B SaaS	5–20 direct, plus the regulators that matter	A small strategy or PMM function, no CI team	India (proposed)

PRIMARY USER

“The keeper of the competitor spreadsheet”

A product marketer, strategy associate or CI analyst. Wants no surprises, and “what changed?” answered in a minute, with evidence.

BUYER

Head of Strategy, VP Product Marketing or CPO

Or the founder, in smaller companies. Wants systematic coverage and decisions they can defend.

THEN EXPAND TO

- Enterprise CI teams in pharma, automotive and manufacturing
- Compliance teams scanning the regulatory horizon
- Investors watching portfolio companies and their competitors

Product, Sales & BD, Compliance and Leadership receive what is routed to them and configure nothing. The engine works in any industry: the planner discovers a domain model per company — EVs and plants for Tata Motors; trials and approvals for Pfizer.

Others do part of the job. SignalLens does it end to end.

	Page-change monitors e.g. Visualping, Distill	Alerts & media monitoring e.g. Google Alerts, Meltwater	CI platforms e.g. Crayon, Klue	Research & company data e.g. Crunchbase, AlphaSense	General AI assistants e.g. ChatGPT, Perplexity	SignalLens
GREAT AT	Simple, reliable detection on almost any page	Broad news coverage, PR and brand analytics	Battlecards and sales enablement	Deep datasets and research search	Fast answers and cited research, on demand	What changed, is it true, why it matters, who needs it
Watches continuously, unasked	●	●	●	◐	○	● on a schedule
Removes noise before alerting	◐	◐	◐	—	—	● four tiers
One alert per real-world event	—	○	◐	●	—	● clustering
Verified quotes, rule-based status	○	○	○	◐	◐	● computed
Versioned history of facts	◐	◐	◐	●	○	● + a year of backfill
Why it matters to <i>your</i> company	○	○	◐	○	◐	● your profile
Routes to the owning team by urgency	◐	◐	●	◐	○	● now or digest

● Yes ◐ Partly ○ No, or typically not — Not applicable

Category-level and qualitative; individual products vary. Research platforms' history and curation cover their own datasets.

Charge for what is watched. Costs follow changes, not checks.

PRICING HYPOTHESES

To validate — no pricing tested yet

H1 Charge for what is watched, not for who reads

Price by monitored subjects, their sources and check frequency.
Unlimited readers, so intelligence spreads across teams.

H2 Land with one team, then expand

Start with strategy or product marketing on 3–5 competitors; add
Compliance, Sales & BD and more subjects.

H3 A short, self-serve trial

A year of history on day one shows value before any live change happens.

H4 Anchored on analyst time saved

And on existing CI and research budgets. Possible tiers: Team,
Business, Enterprise.

UNIT ECONOMICS

≈ **\$4–6** model cost per workspace
per month

Estimate, not a measurement

Estimated from code paths and assumed volumes: one subject plus
regulators, 10 sources, about 15 investigations a month. One-time
onboarding, including a 12-month backfill, is estimated at well under \$1.

Checks	Most end at a hash or “not modified” — no model
Changes	A fast, cheap model, only on what changed
Material events	Reasoning model, with a hard budget per run
Cards	One strong-model call each

Every run records its tokens and cost, so real unit costs come from
measurement, not guesses.

The loop works today. Next: wider coverage, enterprise readiness, measured quality.

NOW v1, built

- The full loop, from plan to learned rules
- Web app with every product screen
- In-app inbox, Slack and a daily digest
- Approval queue for external actions
- Demo lab for trials and training

NEXT

- **JavaScript page rendering** — such pages are detected and skipped today
- **Email digest** and Microsoft Teams
- **Single sign-on** and SCIM
- **More approval-gated actions**, such as CRM updates
- **Evaluation harness** — replay labelled archive history to measure precision and recall
- **Ask the world state** — answers that cite versions and evidence

LATER direction, not commitment

- Push verified cards into CI platforms, CRMs and wikis
- Registries as sources: clinical trials, patents
- Portfolio views for investors

Deliberately never: six chatty agents, blanket crawling of whole sites, or a generic chat box.

We're looking for design partners and pilot customers

WHO

B2B teams in **payments, fintech or B2B SaaS** that track 5–20 competitors and their regulators, without a dedicated CI team.

THE PILOT

Approve a plan for **3–5 competitors** and your regulators. See a year of history on day one, then live cards routed to your teams.

WHAT WE MEASURE TOGETHER

Alert precision, time to detection and hours saved — with targets set from your own baseline.

WHAT YOU BRING

A named owner, **feedback on every card**, and API keys for the model and search providers you prefer.



Know what changed, whether it's true, and why it matters to you.